




Contain Yourself: Staying Undetected Using the Windows Containers Isolation Framework

DEF CON 31 2023

About Me

 [@daniel_Avinoam](https://twitter.com/daniel_Avinoam)

Security Researcher @ 

Interested in Windows internals, reverse engineering and low-level programming

Motorsport fan and weightlifter



Today's Agenda


01



Overview

Windows containers
basics


02



wcifs.sys

Analyzing the main filter
driver responsible for
containers FS isolation

03



Utilization

Bypassing security
products using our
findings

04



Summary

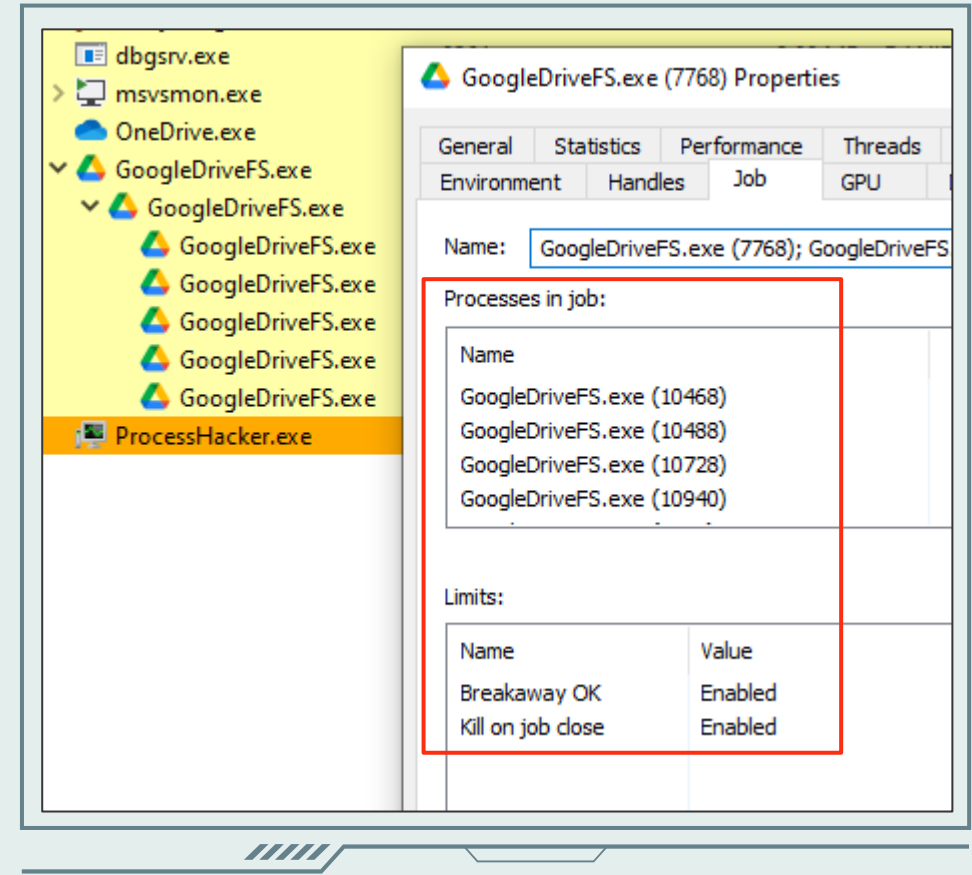
Summary, mitigation
and future research

Windows Containers – The Basics



Jobs

- Objects designed to group several processes and manage them as one unit
- Allows limitation of CPU usage, I/O bandwidth, virtual memory usage, and network activity of managed processes
- Often used by multi-processed applications



Windows Containers – The Basics



Silos

- Like traditional jobs, these objects are used for process grouping with additional features
- Server silo - provides separation of various system resources like the registry, networking, and the object manager
- The Windows kernel detects processes assigned to silos using APIs like *PsIsCurrentThreadInServerSilo* and *PsIsProcessInSilo*

```
NTSTATUS IopUnloadDriver(...)  
{  
    [snip]  
    if (PsIsCurrentThreadInServerSilo())  
    {  
        DbgPrint("Server Silo attempted to unload driver");  
        return STATUS_PRIVILEGE_NOT_HELD;  
    }  
    ...  
}
```



Windows Containers – The Basics



File System Redirection Using Reparse Points

- Reparse points are MFT attributes that can be given to files or directories
- Stores a reparse tag and user-defined data

```
// This header is universal for all points  
struct REPARSE_DATA_BUFFER {  
    /*0*/ ULONG   ReparseTag;  
    /*4*/ USHORT  ReparseDataLength;  
    /*6*/ USHORT  Reserved;  
    /*8*/ UCHAR   DataBuffer[1];  
};
```

- Can be set using `DeviceIoControl` + `FSCTL_SET_REPARSE_POINT` (WRITE primitive needed)





Windows Containers – The Basics

File System Redirection Using Reparse Points

- When a file with a reparse point is opened, it is handled by a file system mini-filter driver according to its reparse tag
- An example of a common tag is `IO_REPARSE_TAG_SYMLINK`, which is how symbolic links work behind the scenes.

<code>IO_REPARSE_TAG_SYMLINK</code>	Used for symbolic link support. See section 2.1.2.4 .
<code>0xA000000C</code>	

https://learn.microsoft.com/en-us/openspecs/windows_protocols/ms-fscc/c8e77b37-3909-4fe6-a4ea-2b9d423b1ee4

```
c:\>"C:\Users\Daniel\Desktop\SysinternalsSuite\junction64.exe" "C:\Users\All Users"

Junction v1.07 - Creates and lists directory links
Copyright (C) 2005-2016 Mark Russinovich
Sysinternals - www.sysinternals.com

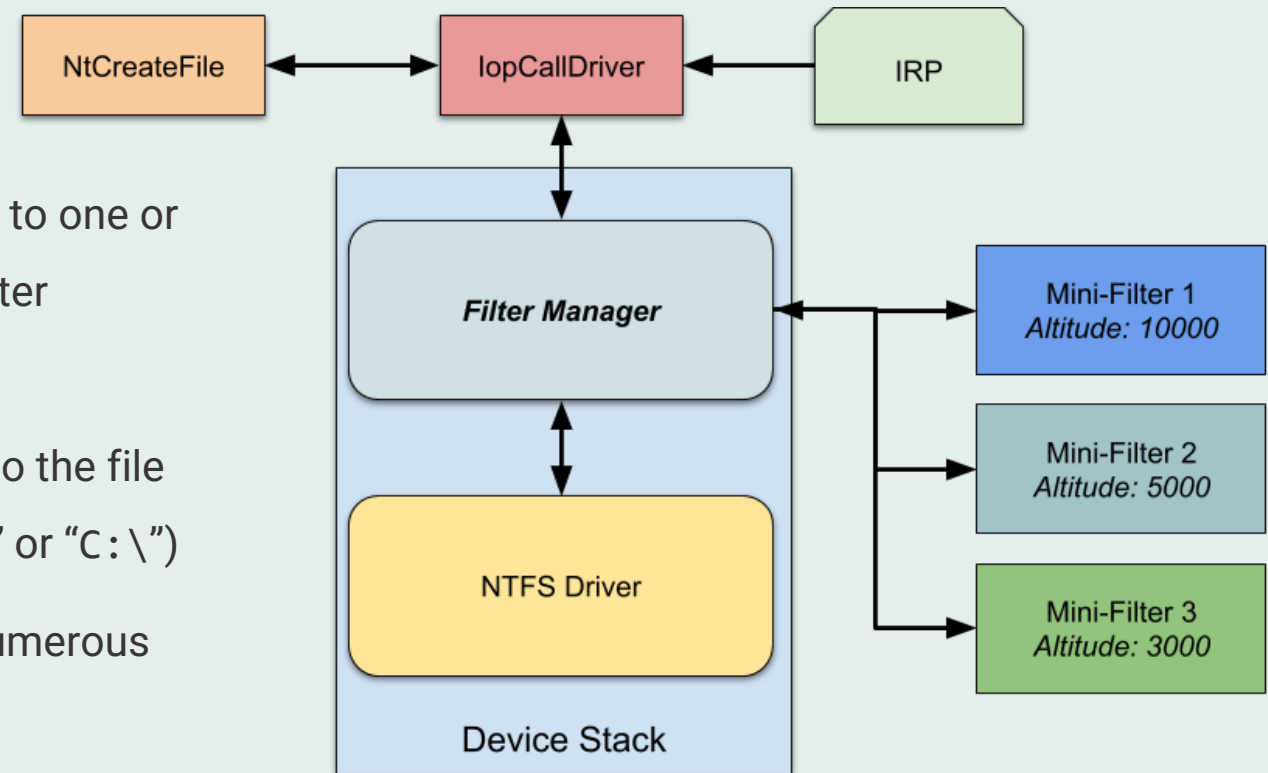
C:\Users\All Users: SYMBOLIC LINK
Print Name       : C:\ProgramData
Substitute Name: \??\C:\ProgramData
```

Windows Containers – The Basics

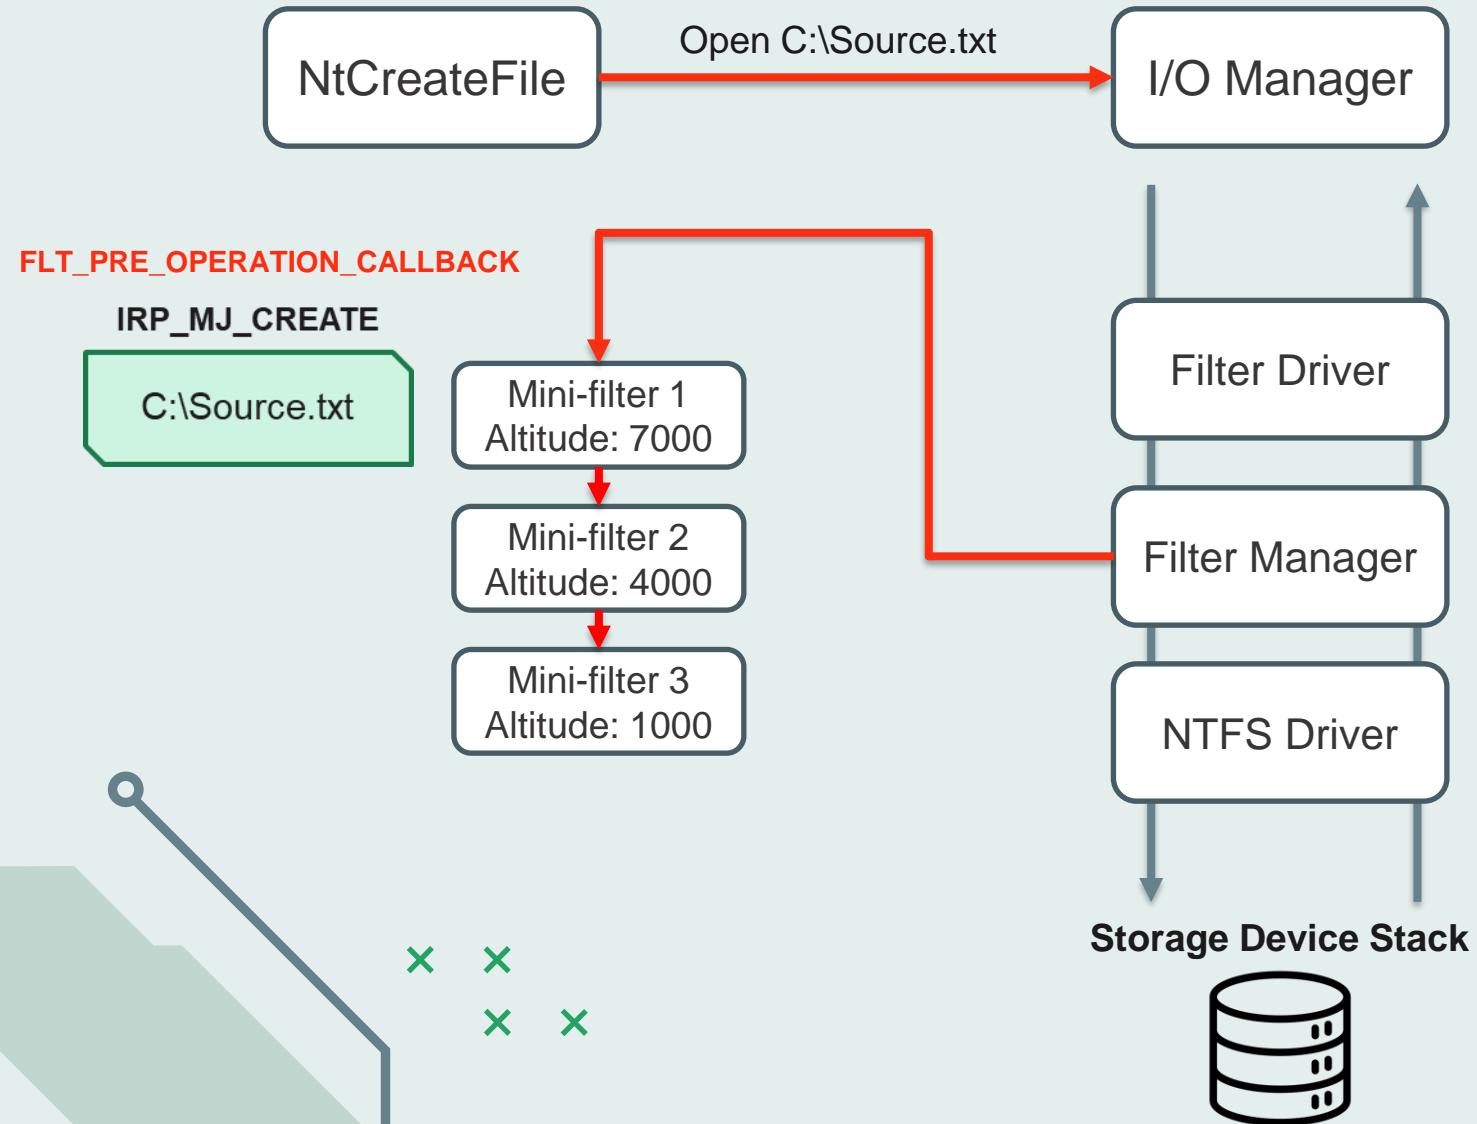


Mini-filters Background

- Filter manager - a legacy filter that manages other “mini” filter drivers (exposes the **Flt** API)
- Each mini-filter can be attached by the manager to one or more volumes, creating what is called a “mini-filter instance”
- Volume = logical storage unit that is presented to the file system as a disk (“\Device\HarddiskVolume3” or “C:\”)
- Can intercept the PRE and POST operations of numerous I/O functions
- Attached and ordered according to an altitude



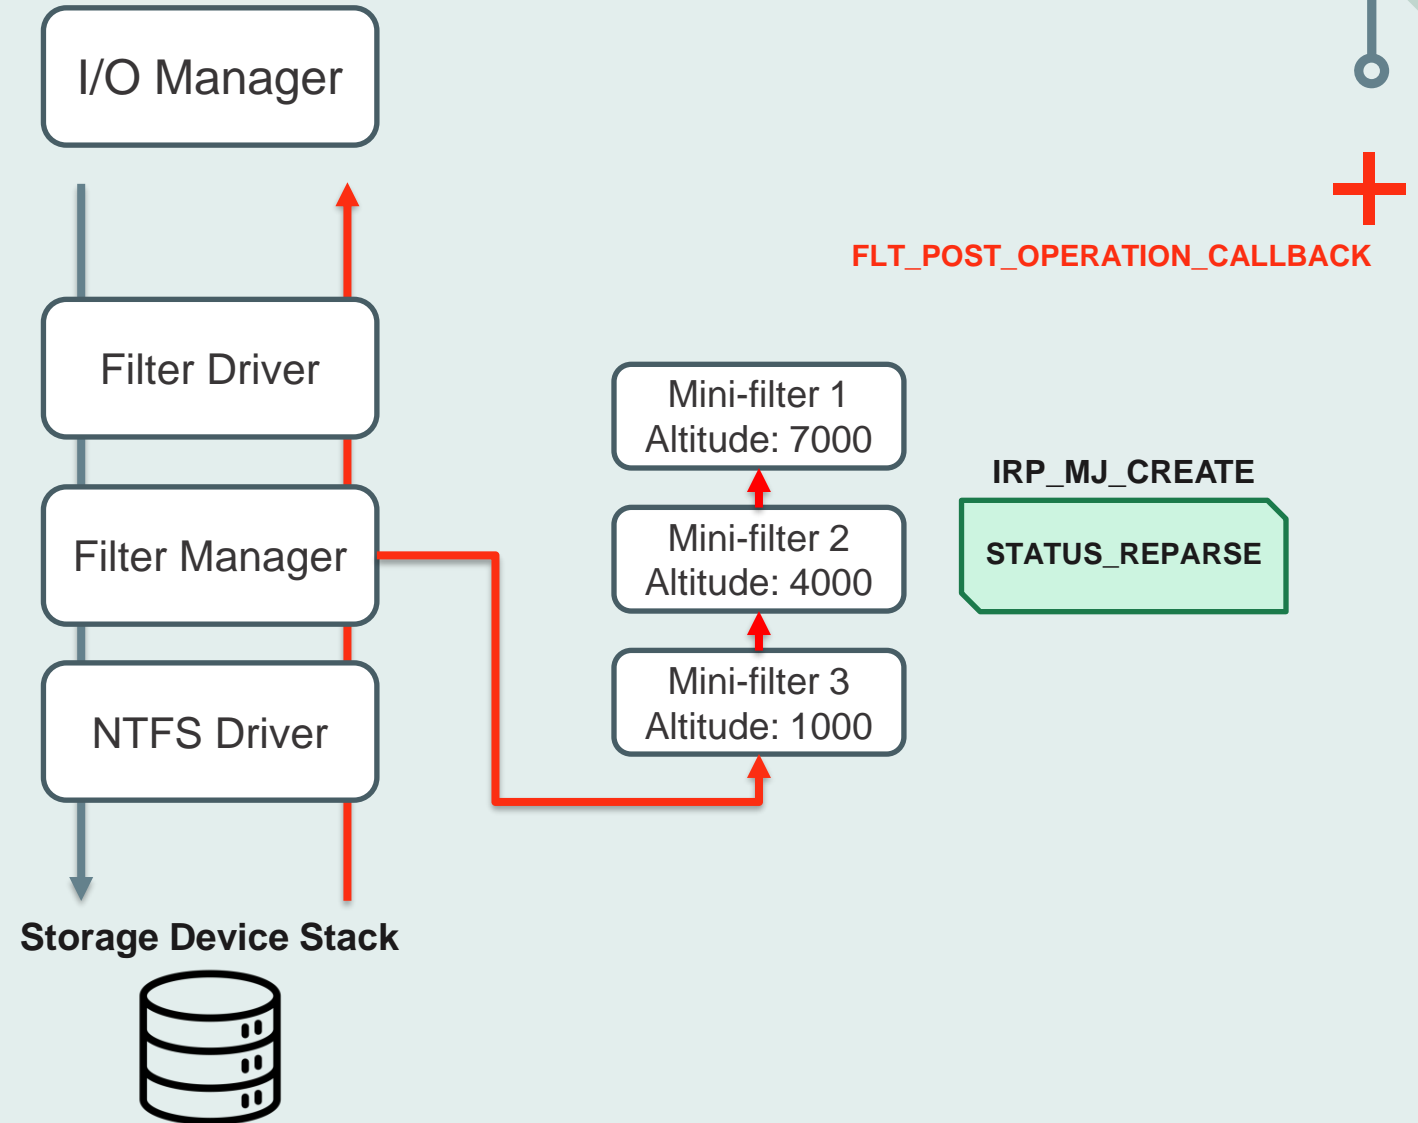
Mini-filters and Reparse Points



Mini-filters and Reparse Points

Option 1:

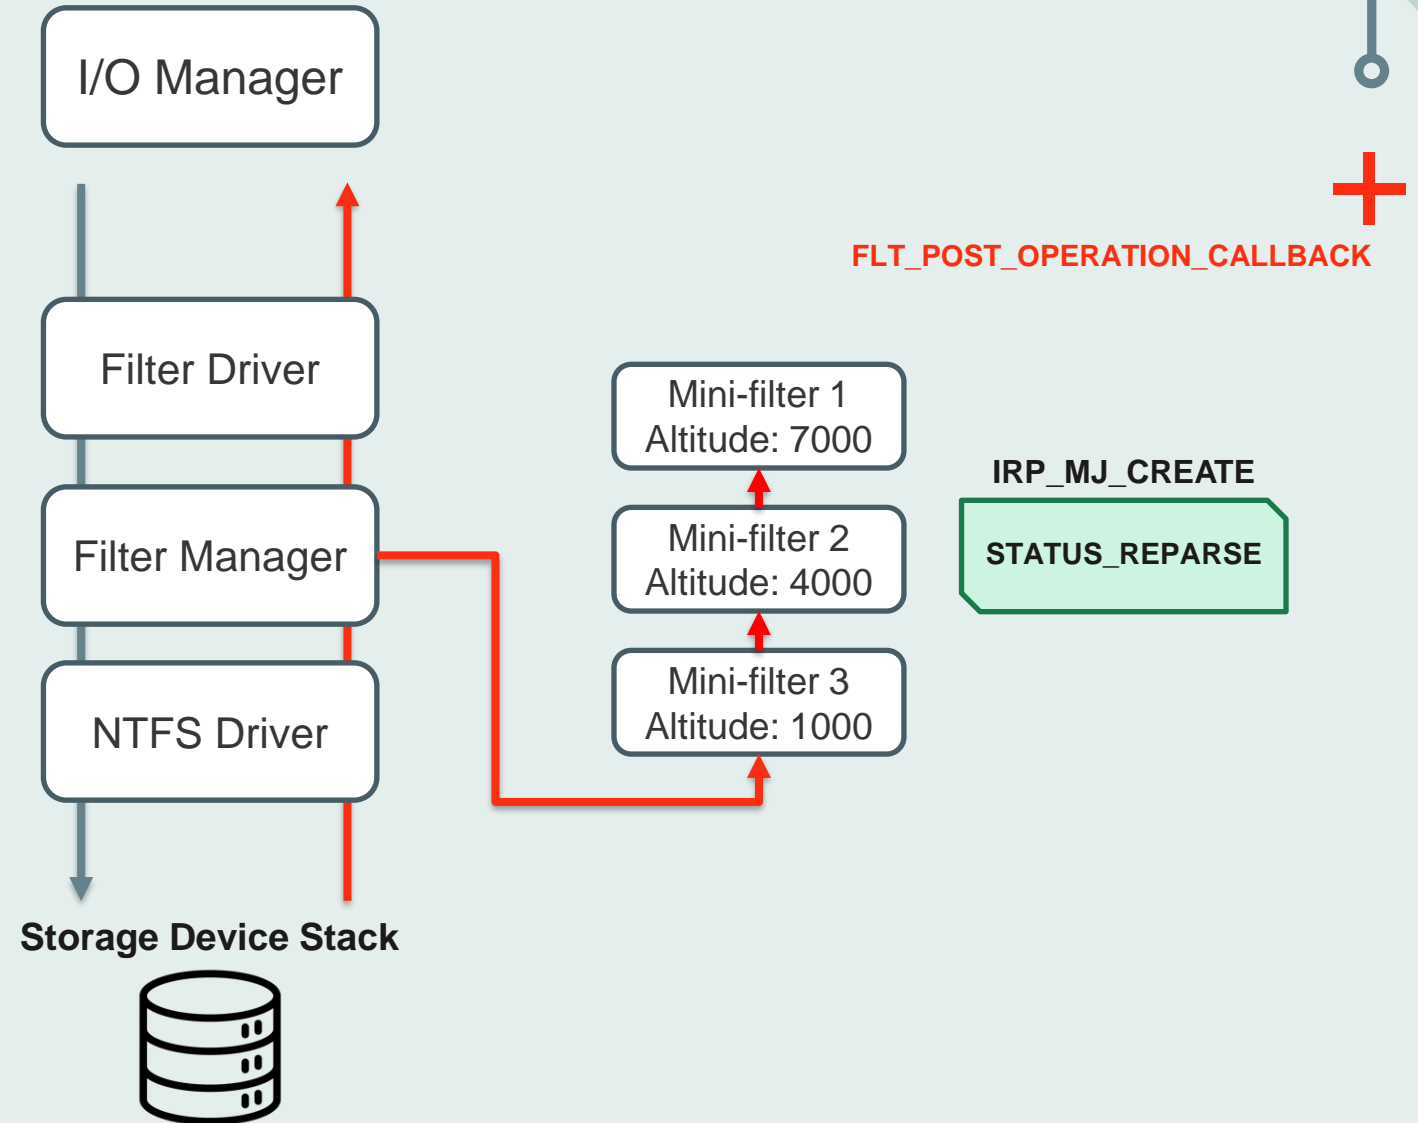
```
FltFsControlFile(...,  
...,  
FSCTL_GET_REPARSE_POINT,  
&ReparsePointData);  
  
ReparsePointData =  
{  
    MINIFILTER_1_TAG;  
    Minifilter1ReparseDataLen;  
    Minifilter1ReparseData[];  
}
```



Mini-filters and Reparse Points

Option 2:

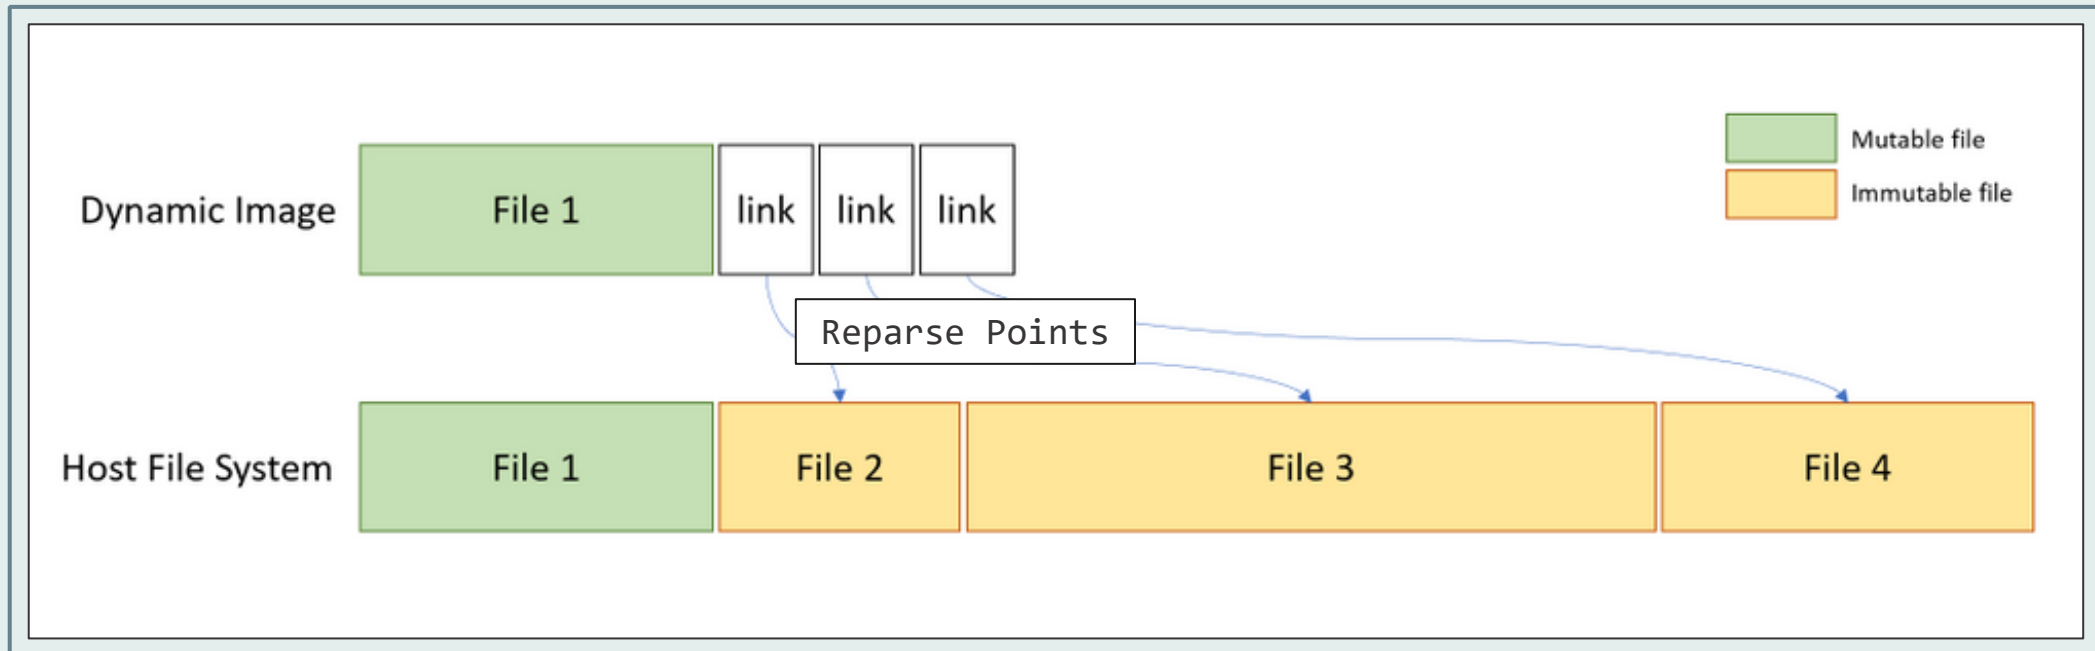
```
FltFsControlFile(...,  
...,  
FSCTL_GET_REPARSE_POINT,  
&ReparsePointData);  
  
ReparsePointData =  
{  
  MINIFILTER_2_TAG;  
  MiniFilter2ReparseDataLen;  
  MiniFilter2ReparseData;  
  {  
    TargetPath = "C:\\\\Target.txt";  
  };  
};
```



Containers File System Separation



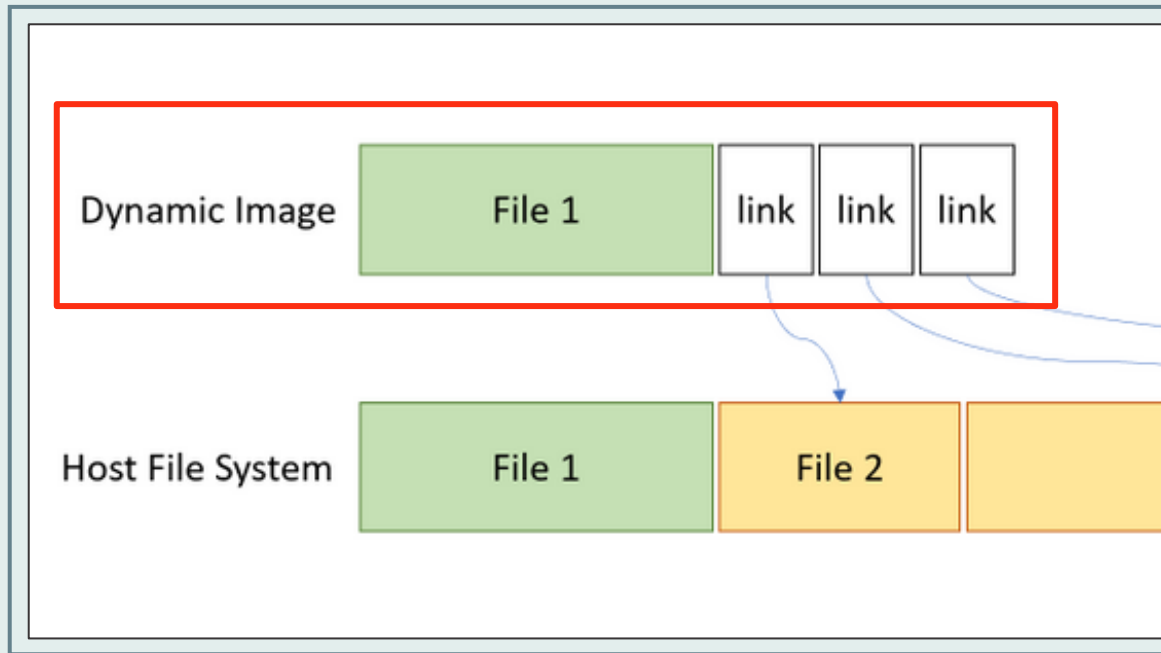
In order to avoid an additional copy of OS files, each container is using a dynamically generated image which points to the original using reparse points



Containers File System Separation



The result is images that contain “ghost files”, which store no actual data but point to a different volume on the system



The screenshot shows a Windows Explorer window with a folder listing on the left and a 'Windows Properties' dialog box on the right. The dialog box is open to the 'General' tab.

Name	Date modified	Type
EFI	8/8/2020 3:38 AM	File fo
OfficePackagesForWDAG	8/8/2020 5:34 AM	File fo
PerfLogs	12/7/2019 11:14 AM	File fo
Program Files	8/8/2020 5:22 AM	File fo
Program Files (x86)		
ProgramData		
Users		
Windows		

Windows Properties

General | Sharing | Security | Previous Versions | Cus

Type: File folder

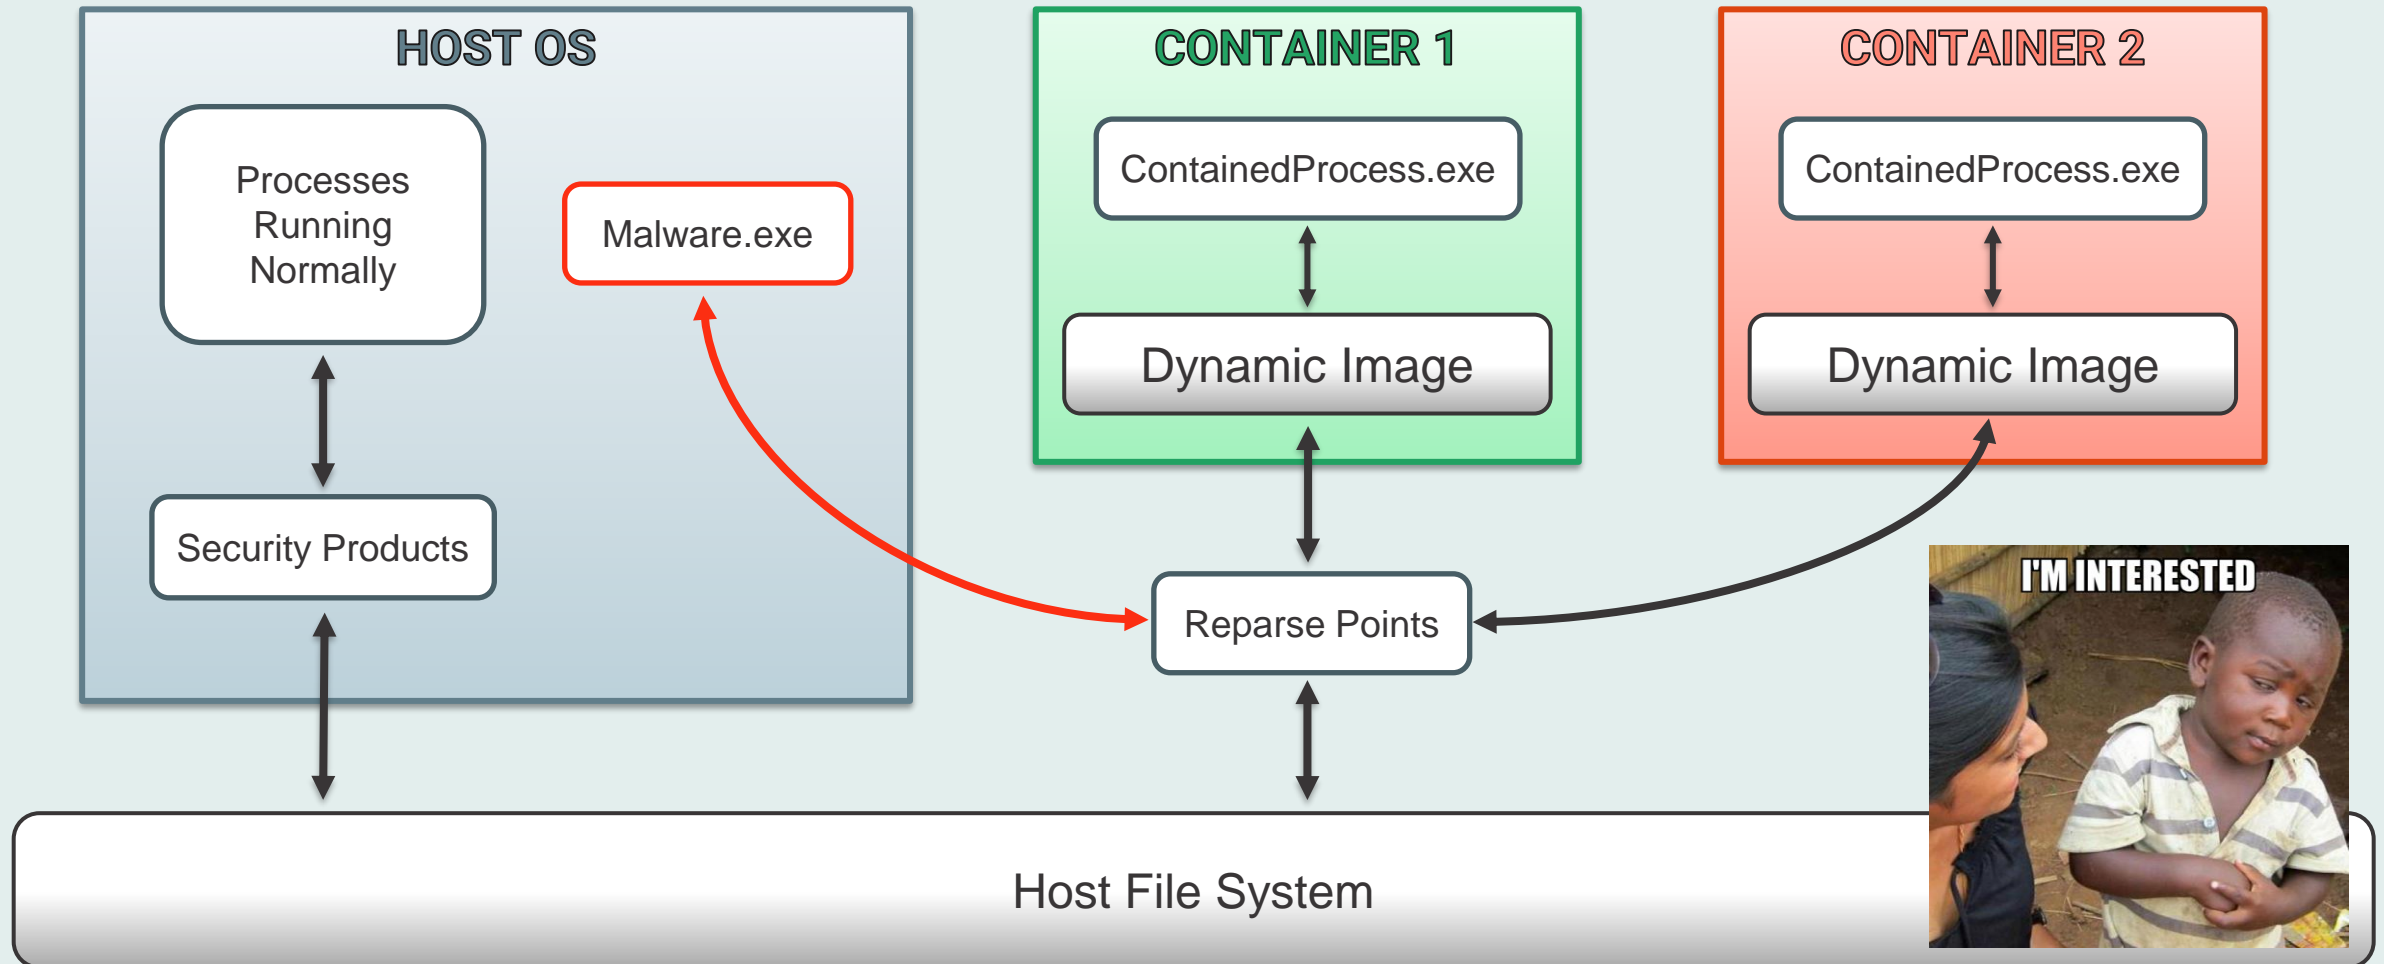
Location: E:\

Size: 1.25 GB (1,348,553,450 bytes)

Size on disk: 57.6 MB (60,432,384 bytes)

Contains: 3,259 Files, 482 Folders

Containers File System Separation




01



Overview

Windows containers
basics

02



wcifs.sys

Analyzing the main filter
driver responsible for
containers FS isolation

03



Utilization

Bypassing security
products using our
findings

04



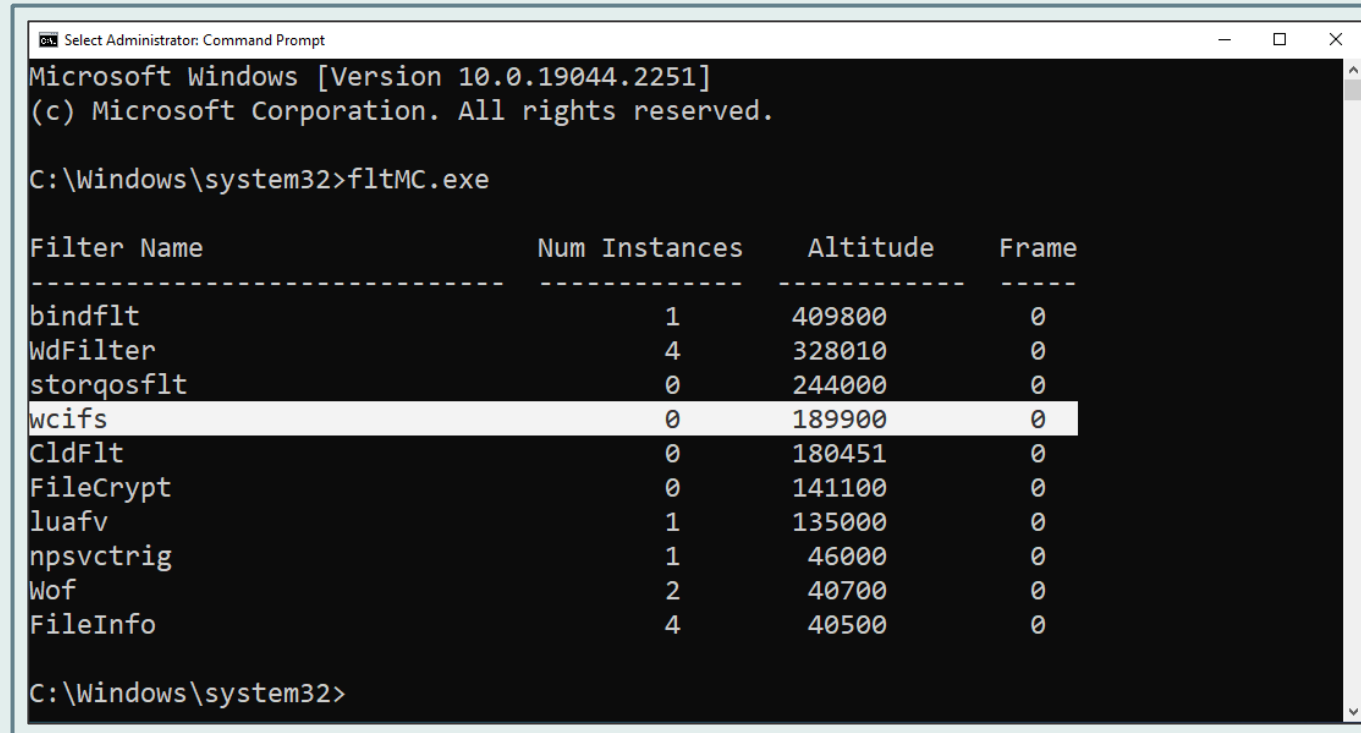
Summary

Summary, mitigation
and future research

wcifs.sys

The Windows Container Isolation FS (wcifs) mini-filter driver is responsible for the file system separation between windows containers and their host

The driver is loaded by default on every Windows system starting from Windows 10



```
Select Administrator: Command Prompt
Microsoft Windows [Version 10.0.19044.2251]
(c) Microsoft Corporation. All rights reserved.

C:\Windows\system32>fltMC.exe

Filter Name                Num Instances  Altitude  Frame
-----
bindflt                    1             409800    0
WdFilter                   4             328010    0
storqosflt                 0             244000    0
wcifs                      0             189900    0
CldFlt                     0             180451    0
FileCrypt                  0             141100    0
luafv                      1             135000    0
npsvc trig                 1              46000    0
Wof                        2              40700    0
FileInfo                   4              40500    0

C:\Windows\system32>
```


wcifs.sys

The main reparse tags associated with this driver are **IO_REPARSE_TAG_WCI_1** and **IO_REPARSE_TAG_WCI_LINK_1**

IO_REPARSE_TAG_WCI_1 0x90001018	Used by the Windows Container Isolation filter. Server-side interpretation only, not meaningful over the wire.
IO_REPARSE_TAG_WCI_LINK_1 0xA0001027	Used by the Windows Container Isolation filter. Server-side interpretation only, not meaningful over the wire.

wcifs.sys

The main reparse tags associated with this driver are `IO_REPARSE_TAG_WCI_1` and `IO_REPARSE_TAG_WCI_LINK_1`

// This header is universal for all points


```
struct REPARSE_DATA_BUFFER {  
    /*0*/ ULONG   ReparseTag; // IO_REPARSE_TAG_WCI_1 / LINK_1  
    /*4*/ USHORT  ReparseDataLength;  
    /*6*/ USHORT  Reserved;  
    /*8*/ UCHAR   DataBuffer[1];  
};
```

```
struct WcifsReparseDataBuffer {  
    /*0*/  ULONG   Version;  
    /*4*/  ULONG   Reserved;  
    /*8*/  GUID     Guid; //Hardcoded Value  
    /*24*/ USHORT  PathStringLength;  
    /*26*/ wchar_t PathStringBuffer[100];  
};
```

x x
x x

wcifs.sys – Reverse Engineering

All there is left to do is to attach wcifs to a volume, place a breakpoint on its POST_OP callback while debugging it and see how its reparse points are being handled



wcifs.sys – Reverse Engineering

For the POST_OP callback to invoke, either `FLT_PREOP_SUCCESS_WITH_CALLBACK` or `FLT_PREOP_SYNCHRONIZE` must be returned in the PRE_OP!

```
FLT_PREOP_CALLBACK_STATUS WcPreCreate(...)
{
    [snip]

    if (!WcUnionsExistForInstance(FltObjects->Instance, ...))
    {
        return FLT_PREOP_SUCCESS_NO_CALLBACK;
    }
    ...

    return FLT_PREOP_SYNCHRONIZE;
}
```

```
BOOL WcUnionsExistForInstance(...)
{
    [snip]

    Silo = IoGetSilo(FileObject);
    IsHostSilo = PsIsHostSilo(Silo);

    if (IsHostSilo)
        return FALSE;

    if (!NT_SUCCESS(PsGetSiloContext(Silo, ...)))
        return FALSE;

    ...
}
```

A context is a structure that is defined by the mini-filter driver and that can be associated with a filter manager object, like files, instances and silos.

× ×

wcifs.sys – Creating a Fake Container

1. Create a silo and insert our process into it
2. Inform the driver our silo is representing a container, so it will create a union context and handle it accordingly



wcifs.sys – Creating a Fake Container

1. **Create a silo and insert our process into it**
2. Inform the driver our silo is representing a container, so it will create a union context and handle it accordingly

```
void CreateSilo(...)
{
    // Create a job
    SECURITY_ATTRIBUTES securityAttributes {};
    HANDLE jobHandle = CreateJobObjectA(&securityAttributes, "ContainYourselfJob");

    // Convert to a silo
    SetInformationJobObject(jobHandle, JobObjectCreateSilo, nullptr, 0);

    // Assign our process
    AssignProcessToJobObject(jobHandle, GetCurrentProcess());
}
```

wcifs.sys – Creating a Fake Container

1. Create a silo and insert our process into it
2. **Inform the driver our silo is representing a container, so it will create a union context and handle it accordingly**

```
void CreateSilo(...)
{
    // Create a job
    SECURITY_ATTRIBUTES securityAttributes {};
    HANDLE jobHandle = CreateJobObjectA(&securityAttributes, "ContainYourselfJob");

    // Convert to a silo
    SetInformationJobObject(jobHandle, JobObjectCreateSilo, nullptr, 0);

    // Assign our process
    AssignProcessToJobObject(jobHandle, GetCurrentProcess());
}
```

wcifs.sys – Creating a Fake Container

1. Create a silo and insert our process into it
2. **Inform the driver our silo is representing a container, so it will create a union context and handle it accordingly**

```
struct WcifsPortMessage {
    /*0*/ DWORD   MessageCode = SetUnion; // SetUnion = 1
    /*4*/ DWORD   MessageSize;
    /*8*/ char    MessageData[1];
};

struct WcifsPortMessageSetUnion {
    /*0*/  DWORD   MessageVersionOrCode;
    /*4*/  DWORD   MessageSize;
    /*8*/  DWORD   NumberOfUnions;
    /*12*/ wchar_t InstanceName[50];
    /*112*/ DWORD   InstanceNameLength;
    /*116*/ DWORD   ReparseTag;
    /*120*/ DWORD   ReparseTagLink;
    /*124*/ DWORD   Unknown;
    /*128*/ HANDLE   SiloHandle;
    /*136*/ char    UnionData[];
};

struct VolumeUnion {
    /*0*/  GUID   Guid; // Hardcoded value
    /*16*/ BOOL   IsSourceVolume;
    /*20*/ DWORD   OffsetOfVolumeName;
    /*24*/ WORD    SizeOfVolumeName;
    /*26*/ WORD    GuidFlags;
};

struct ContainerRootId {
    /*0*/  USHORT   Size;
    /*2*/  USHORT   Length;
    /*4*/  USHORT   MaximumLength;
    /*6*/  wchar_t  Buffer[23];
};
```

The diagram illustrates the relationship between the structures. An arrow points from the `MessageData` field in `WcifsPortMessage` to the `WcifsPortMessageSetUnion` structure. Another arrow points from the `UnionData` field in `WcifsPortMessageSetUnion` to the `VolumeUnion` structure. A third arrow points from the `OffsetOfVolumeName` field in `VolumeUnion` to the `ContainerRootId` structure.

wcifs.sys – Creating a Fake Container

1. Create a silo and insert our process into it
2. **Inform the driver our silo is representing a container, so it will create a union context and handle it accordingly**

```
struct WcifsPortMessage {
    DWORD  MessageCode = SetUnion; // SetUnion = 1
    DWORD  MessageSize = sizeof(WcifsPortMessage);
    WcifsPortMessageSetUnion Message;
};


struct WcifsPortMessageSetUnion {
    DWORD  MessageVersionOrCode = 1;
    DWORD  MessageSize = sizeof(WcifsPortMessageSetUnion);
    DWORD  NumberOfUnions = 2;
    wchar_t InstanceName[50] = L"wcifs Instance";
    DWORD  InstanceNameLength;
    DWORD  ReparseTag = IO_REPARSE_TAG_WCI_1;
    DWORD  ReparseTagLink = IO_REPARSE_TAG_WCI_LINK_1;
    DWORD  Unknown;
    HANDLE SiloHandle;
    VolumeUnion SourceVolumeUnion;
    VolumeUnion TargetVolumeUnion;
    ContainerRootId SourceVolumeContainerRootId; // "\\Device\\HarddiskVolume5"
    ContainerRootId TargetVolumeContainerRootId; // "\\Device\\HarddiskVolume3"
};
```

wcifs.sys – Processing Reparse Points


IO_REPARSE_TAG_WCI_LINK_1

- As its name suggests, this tag acts as a regular link between two files
- Used when files are opened for read only
- The driver reads the relative path stored in the reparse point and redirects the call to the volume the container directs to using

IoReplaceFileName



```
FltCreateFile("\\Device\\HarddiskVolume3\\dest\\file.txt");
```





```
NtCreateFile("\\Device\\HarddiskVolume5\\source\\file.txt",  
FILE_GENERIC_READ);
```

wcifs.sys

wcifs.sys – Processing Reparse Points

IO_REPARSE_TAG_WCI_1

- File “Expansion”
- Acts as Copy-On-Open protection
- Used when files are opened for write
- Opens a work item that uses `FltReadFile` + `FltWriteFile`
- Source file is deleted when the destination does not exist



```
NtCreateFile("\\Device\HarddiskVolume5\source\file.txt",  
FILE_GENERIC_READ | FILE_GENERIC_WRITE);
```



```
FltWriteFile("\\Device\HarddiskVolume5\source\file.txt");
```

wcifs.sys – Processing Reparse Points

IO_REPARSE_TAG_WCI_1

- Source file is deleted when the destination does not exist

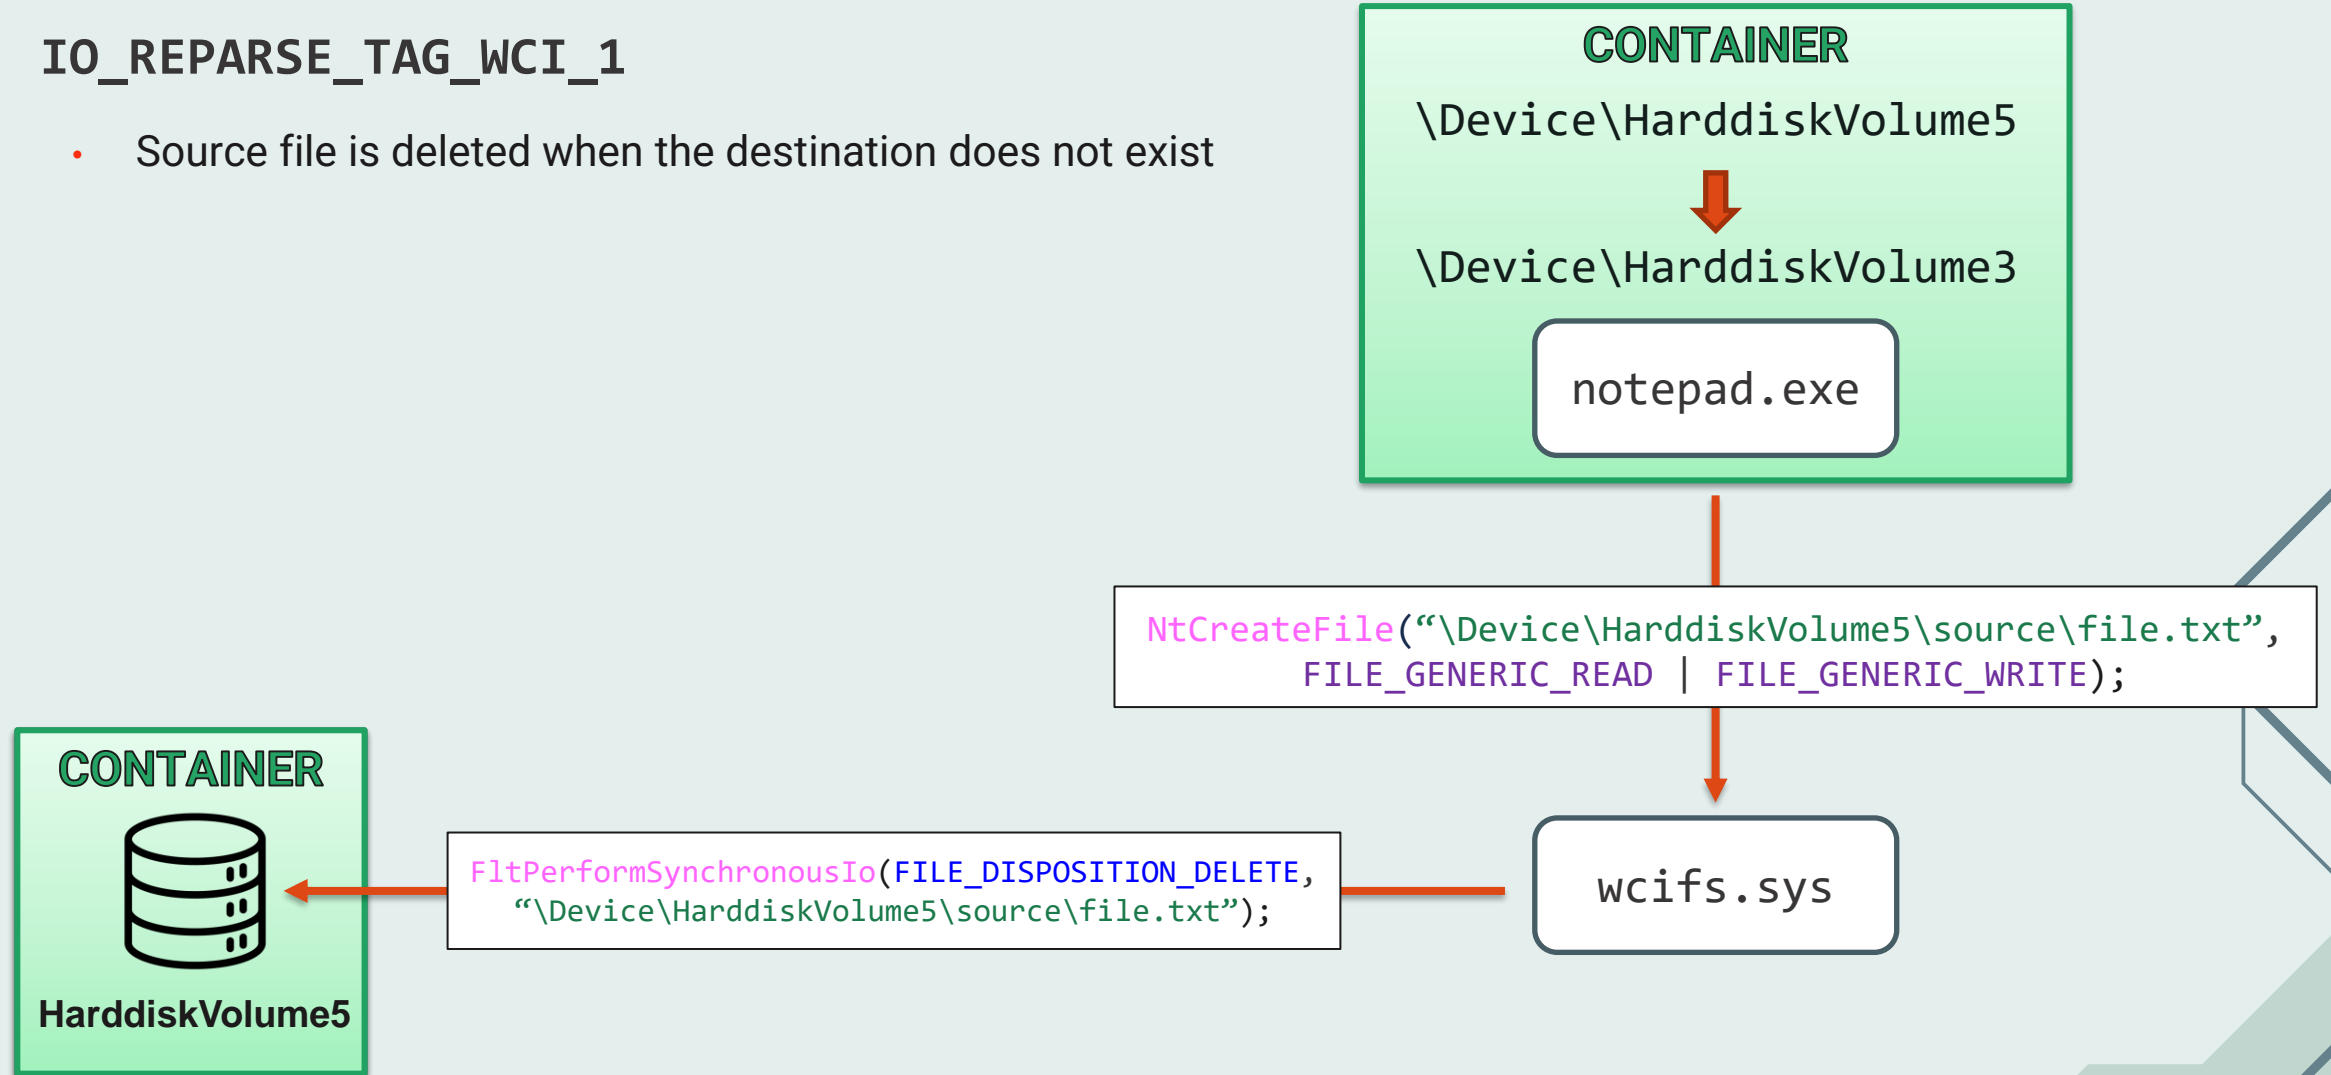
```
status = WcProcessWciReparsePointOpen(CallbackData, FltObjects->Instance, FltObjects->FileObject)
if (status == STATUS_OBJECT_NAME_NOT_FOUND)
{
    NewCallbackData = 0;
    status = FltAllocateCallbackData(FltObjects->Instance, FltObjects->FileObject, &NewCallbackData);
    if (status == STATUS_SUCCESS)
    {
        NewCallbackData->Iopb->MajorFunction = IRP_MJ_SET_INFORMATION;
        NewCallbackData->Iopb->OperationFlags = SL_INFO_IGNORE_READONLY_ATTRIBUTE;
        NewCallbackData->Iopb->Parameters.SetFileInformation.Length = 1;
        NewCallbackData->Iopb->Parameters.SetFileInformation.FileInformationClass = FileDispositionInformation;
        NewCallbackData->Iopb->Parameters.SetFileInformation.ParentOfTarget = 0i64;
        NewCallbackData->Iopb->Parameters.SetFileInformation.ReplaceIfExists = 0i64;
        InfoBuffer = FILE_DISPOSITION_DELETE;
        NewCallbackData->Iopb->Parameters.SetFileInformation.InfoBuffer = &InfoBuffer; //FILE_DISPOSITION_DELETE

        FltPerformSynchronousIo(NewCallbackData);
    }
}
```

wcifs.sys – Processing Reparse Points

IO_REPARSE_TAG_WCI_1

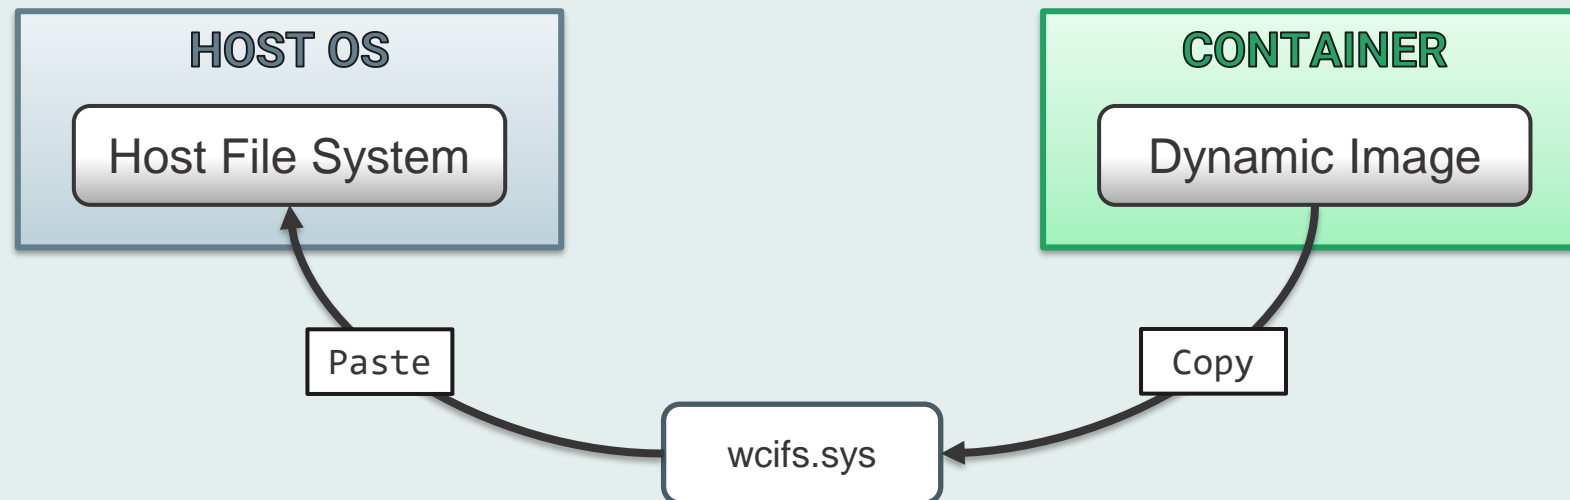
- Source file is deleted when the destination does not exist



wcifs.sys – Other Features

Copy & Paste Files

- Another functionality this driver offers is to copy & paste files, **without the need to enter a container / use reparse points**
- Used when files from needs to be transferred to/from a container's volume
- `FltReadFile` + `FltWriteFile`



wcifs.sys – Other Features

Copy & Paste Files

```
struct WcifsPortMessage {
    /*0*/  DWORD   MessageCode = WcCopyFile; // WcCopyFile = 4
    /*4*/  DWORD   MessageSize = sizeof(WcifsPortMessage);
    /*8*/  WcifsPortMessageCopyFile Message;
};

struct WcifsPortMessageCopyFile {
    /*0*/  DWORD   MessageVersionOrCode = 148;
    /*4*/  DWORD   MessageSize = sizeof(WcifsPortMessageCopyFile);
    /*8*/  wchar_t InstanceName[50] = L"wcifs Instance";
    /*108*/ DWORD   InstanceNameLength;
    /*112*/ DWORD   ReparseTag;
    /*116*/ DWORD   OffsetToSourceContainerRootId;
    /*120*/ DWORD   SizeOfSourceContainerRootId;
    /*124*/ DWORD   OffsetToTargetContainerRootId;
    /*128*/ DWORD   SizeOfTargetContainerRootId;
    /*132*/ DWORD   OffsetToSourceFileRelativePath;
    /*136*/ DWORD   SizeOfSourceFileRelativePath;
    /*140*/ DWORD   OffsetToTargetFileRelativePath;
    /*144*/ DWORD   SizeOfTargetFileRelativePath;
    /*148*/ char    UnionData[]; // 2 * ContainerRootId + source & target relative paths
};
```

wcifs.sys – Summary

So, what have we accomplished so far?

- Create a silo, insert our process into it and register it as a fake container by communicating with the wcifs driver
 1. `IO_REPARSE_TAG_WCI_LINK_1`
 - Open one file and receive the handle of another
 2. `IO_REPARSE_TAG_WCI_1`
 - Override a file with the content of another (`FltReadFile` + `FltWriteFile`)
 - Delete a file (`FltPerformSynchronousIo`)
- Copy & paste a file (`FltReadFile` + `FltWriteFile`)


01



Overview

Windows containers
basics


02



wcifs.sys

Analyzing the main filter
driver responsible for
containers FS isolation

03



Utilization

Bypassing security
products using our
findings

04



Summary

Summary, mitigation
and future research

Security Products Mini-filter Bypass



- Security products use mini-filters to monitor I/O activity
- `FltReadFile`, `FltWriteFile`, `FltPerformSynchronousIo`

“ **[function]** causes the request to be sent to the mini-filter driver instances attached below the initiating instance and to the file system. The specified instance and the Instances attached above it do not receive the request.” - MSDN

<https://learn.microsoft.com/en-us/windows-hardware/drivers/ddi/fltkernel/nf-fltkernel-fltreadfile>

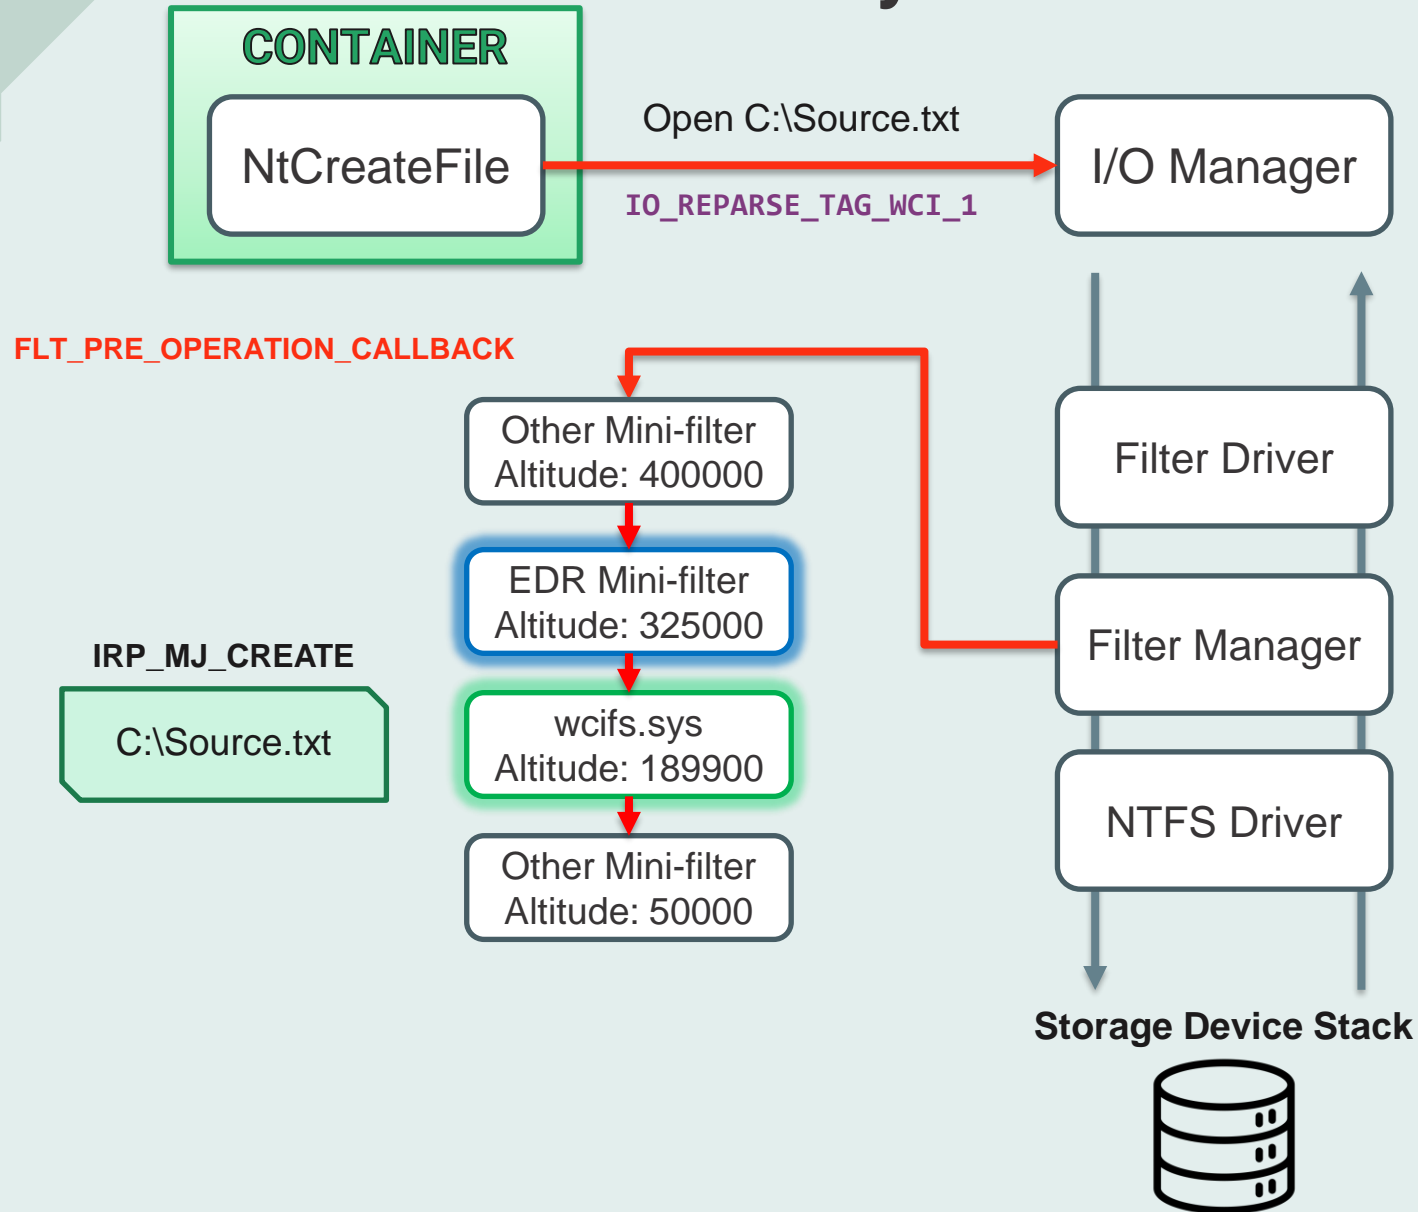
FSFilter Anti-Virus 320000-329999 Includes filter drivers that detect and disinfect viruses during file I/O.

<https://learn.microsoft.com/en-us/windows-hardware/drivers/ifs/load-order-groups-and-altitudes-for-minifilter-drivers>

Filter Name	Num Instances	Altitude	Frame
-----	-----	-----	----
bindflt	1	409800	0
WdFilter	4	328010	0
storqosflt	0	244000	0
wcifs	0	189900	0



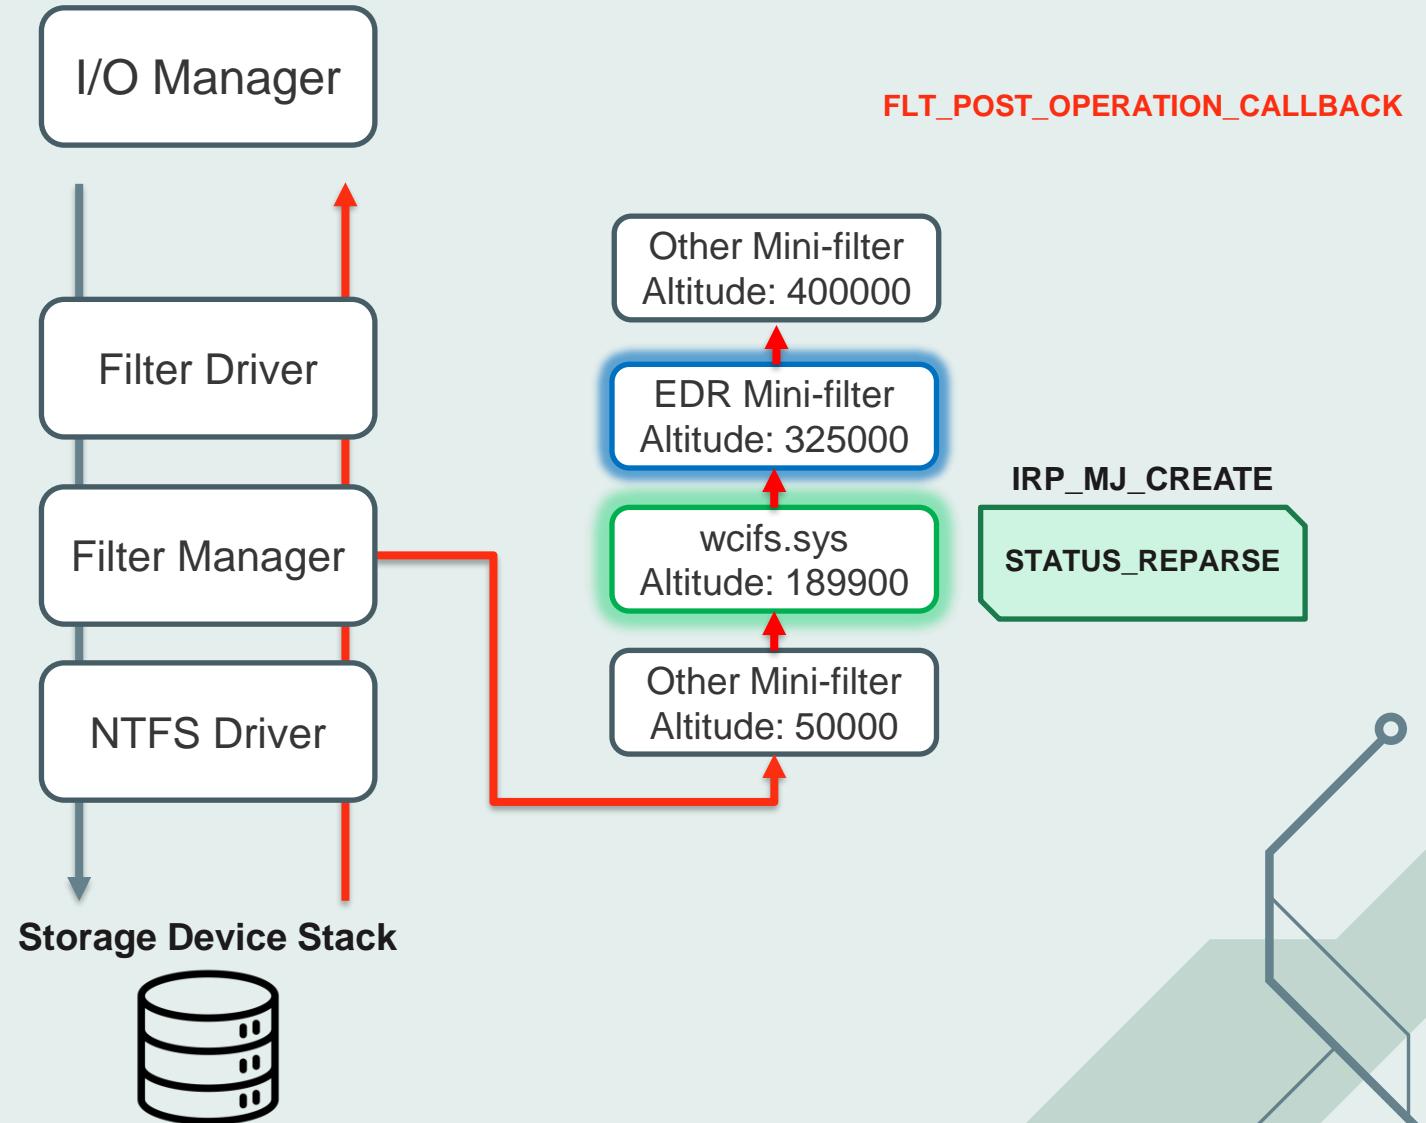
Security Products Mini-filter Bypass



Security Products Mini-filter Bypass



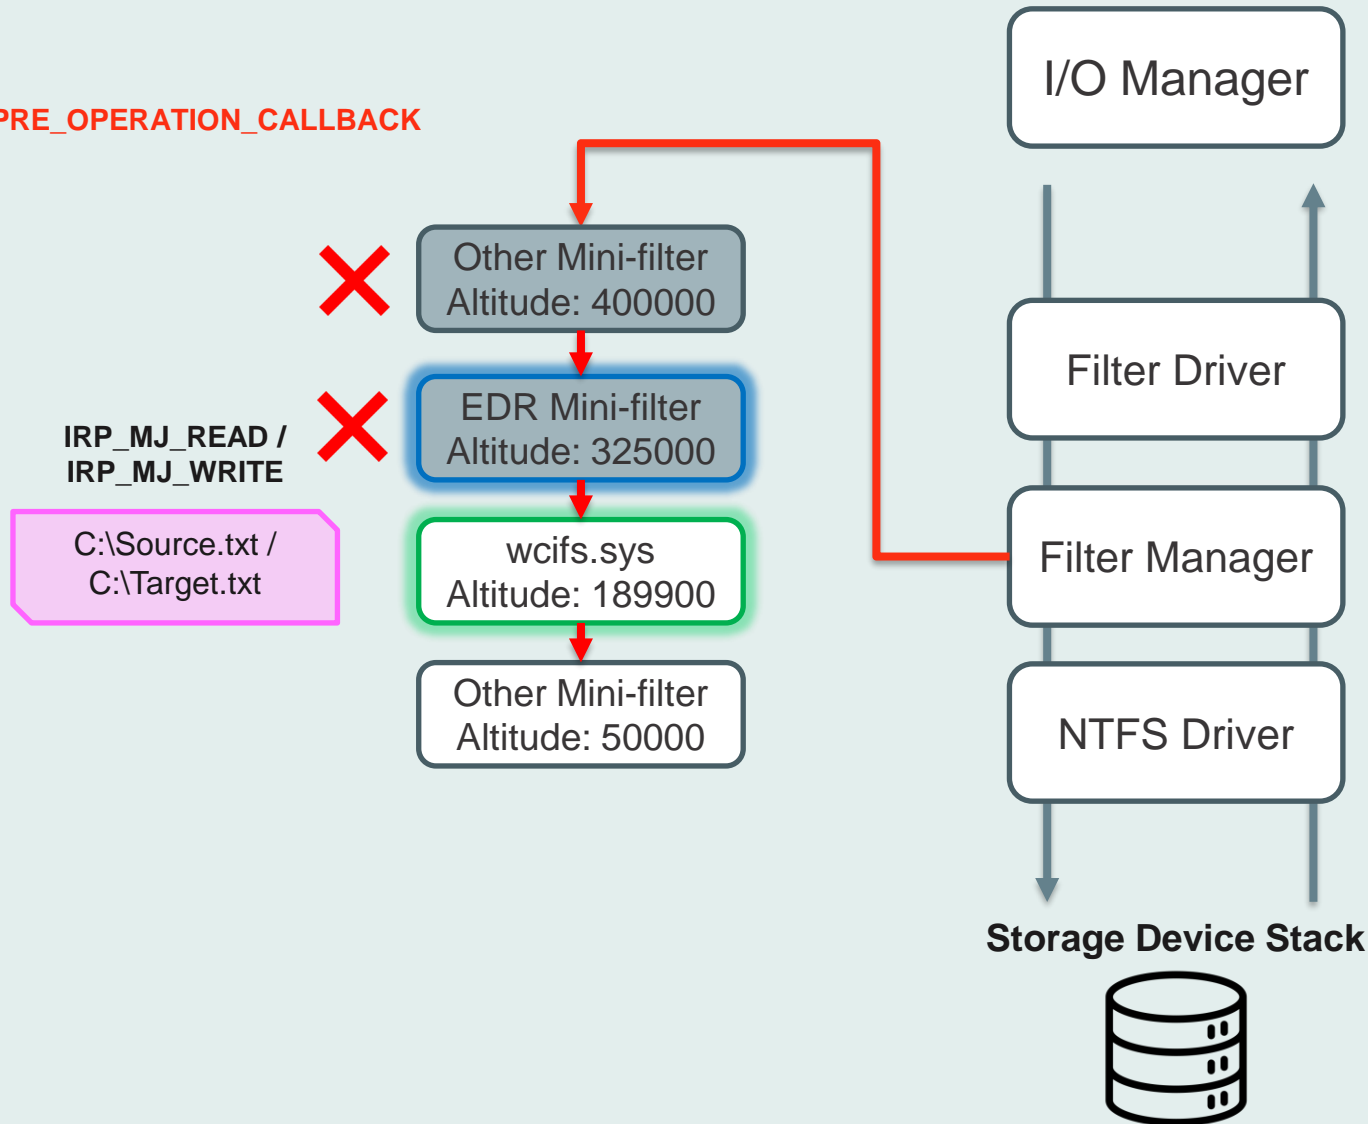
```
ReparsePointData =  
{  
    IO_REPARSE_TAG_WCI_1;  
    MiniFilter2ReparseDataLen;  
    MiniFilter2ReparseData  
    {  
        TargetPath = "C:\\\\Target.txt";  
    };  
}  
  
FltReadFile("C:\\\\Source.txt");  
FltWriteFile("C:\\\\Target.txt");  
OR  
FltPerformSynchronousIo(  
    FILE_DISPOSITION_DELETE,  
    "C:\\\\Source.txt");  
  
return FLT_POSTOP_FINISHED_PROCESSING;
```



Security Products Mini-filter Bypass



FLT_PRE_OPERATION_CALLBACK



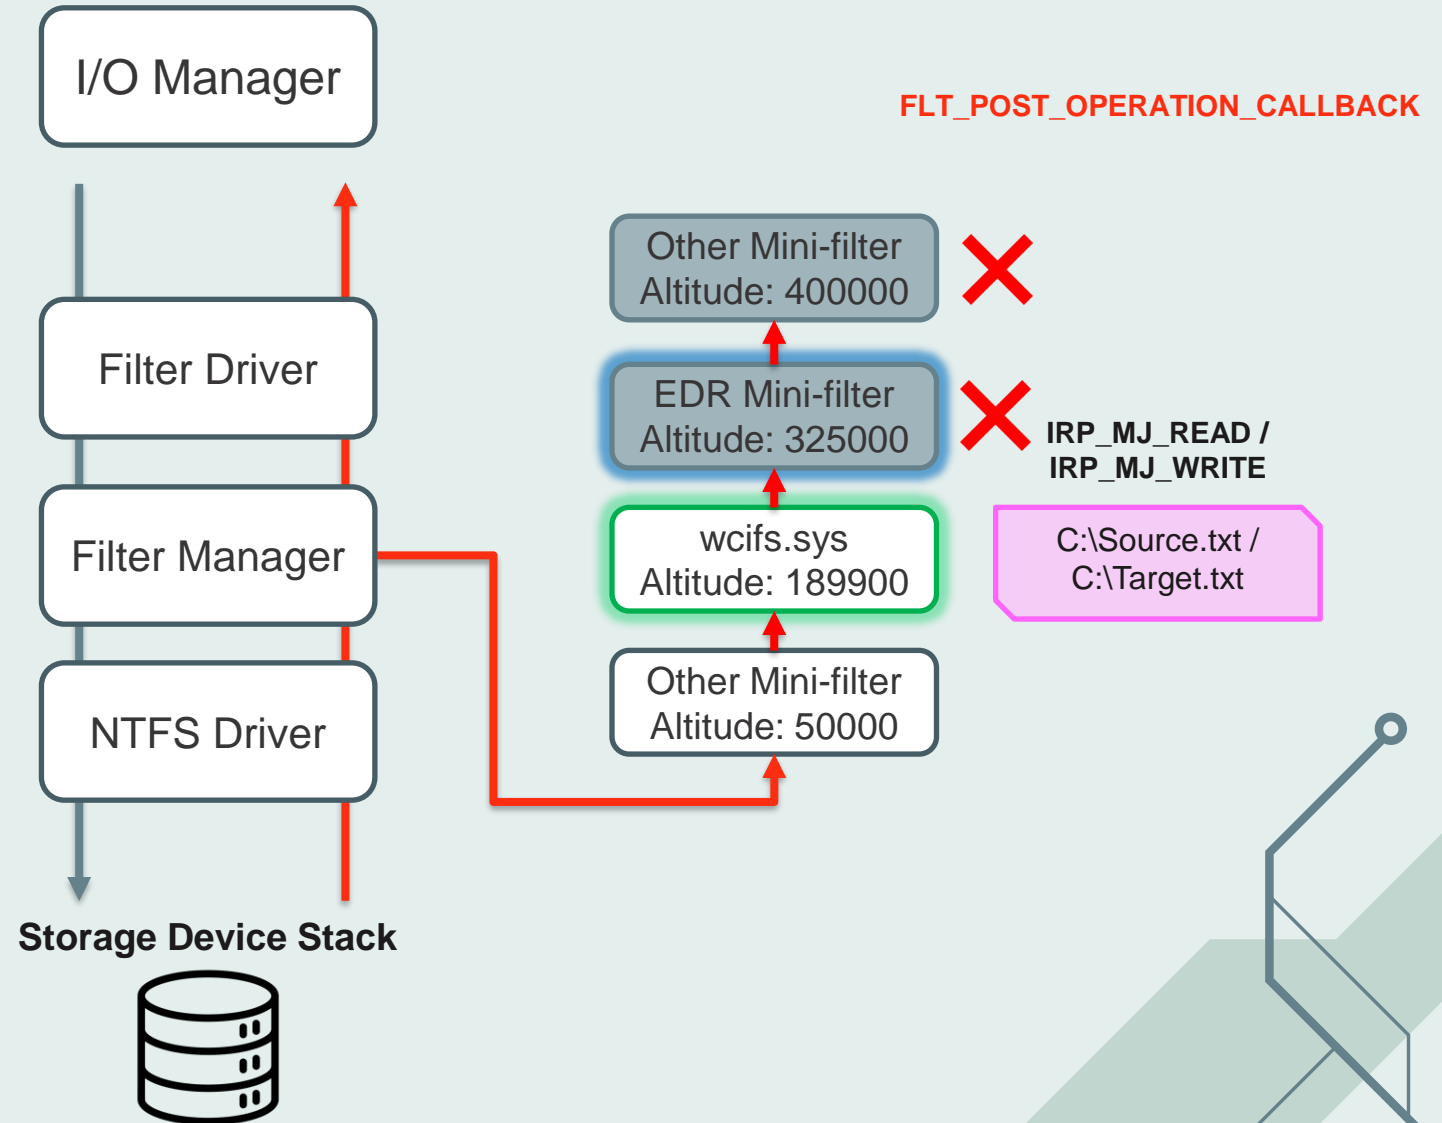
```
FltReadFile("C:\\Source.txt");  
FltWriteFile("C:\\Target.txt");  
OR  
FltPerformSynchronousIo(  
FILE_DISPOSITION_DELETE,  
"C:\\Source.txt");
```

"The specified instance and the Instances attached above it **do not receive** the request."

Security Products Mini-filter Bypass



```
ReparsePointData =  
{  
    IO_REPARSE_TAG_WCI_1;  
    MInifilter2ReparseDataLen;  
    MInifilter2ReparseData  
    {  
        TargetPath = "C:\\\\Target.txt";  
    };  
}  
  
FltReadFile("C:\\\\Source.txt");  
FltWriteFile("C:\\\\Target.txt");  
OR  
FltPerformSynchronousIo(  
    FILE_DISPOSITION_DELETE,  
    "C:\\\\Source.txt");  
return FLT_POSTOP_FINISHED_PROCESSING;
```



Security Products Mini-filter Bypass

EDR 101: Ransomware/Wiper Protection

- Security products employ algorithms that analyze mini-filter I/O logs, searching for specific patterns to detect file system-based malware and prevent them before any irreversible damage is done
- Most EDRs rely on a set of standard principles to categorize a process as ransomware/wiper:
 - I. Process opens handles to a vast number of files
 - II. Process reads data from a file and then writes **to the same file**, making the file's data inaccessible (using pre read/write callbacks)




Security Products Mini-filter Bypass

Creating an Undetectable Wiper

1. Create an empty file that will be our target file. Write a buffer of zeros/random data to it.
2. Traverse each file on the system and for each:
 - I. Set an `IO_REPARSE_TAG_WCI_1` reparse point on the source file that will point to the target file
3. Create a silo, assign the current process to it and register it as a fabricated container to wcifs where both source and target volumes are the main one (`\Device\HarddiskVolume3`).
4. Traverse each file on the system and for each:
 - I. Open the file using `CreateFile` – the files will be overridden with the target file data by the wcifs driver, **the call will not trigger security mini-filter drivers callback function**


Security Products Mini-filter Bypass

Creating an Undetectable Wiper



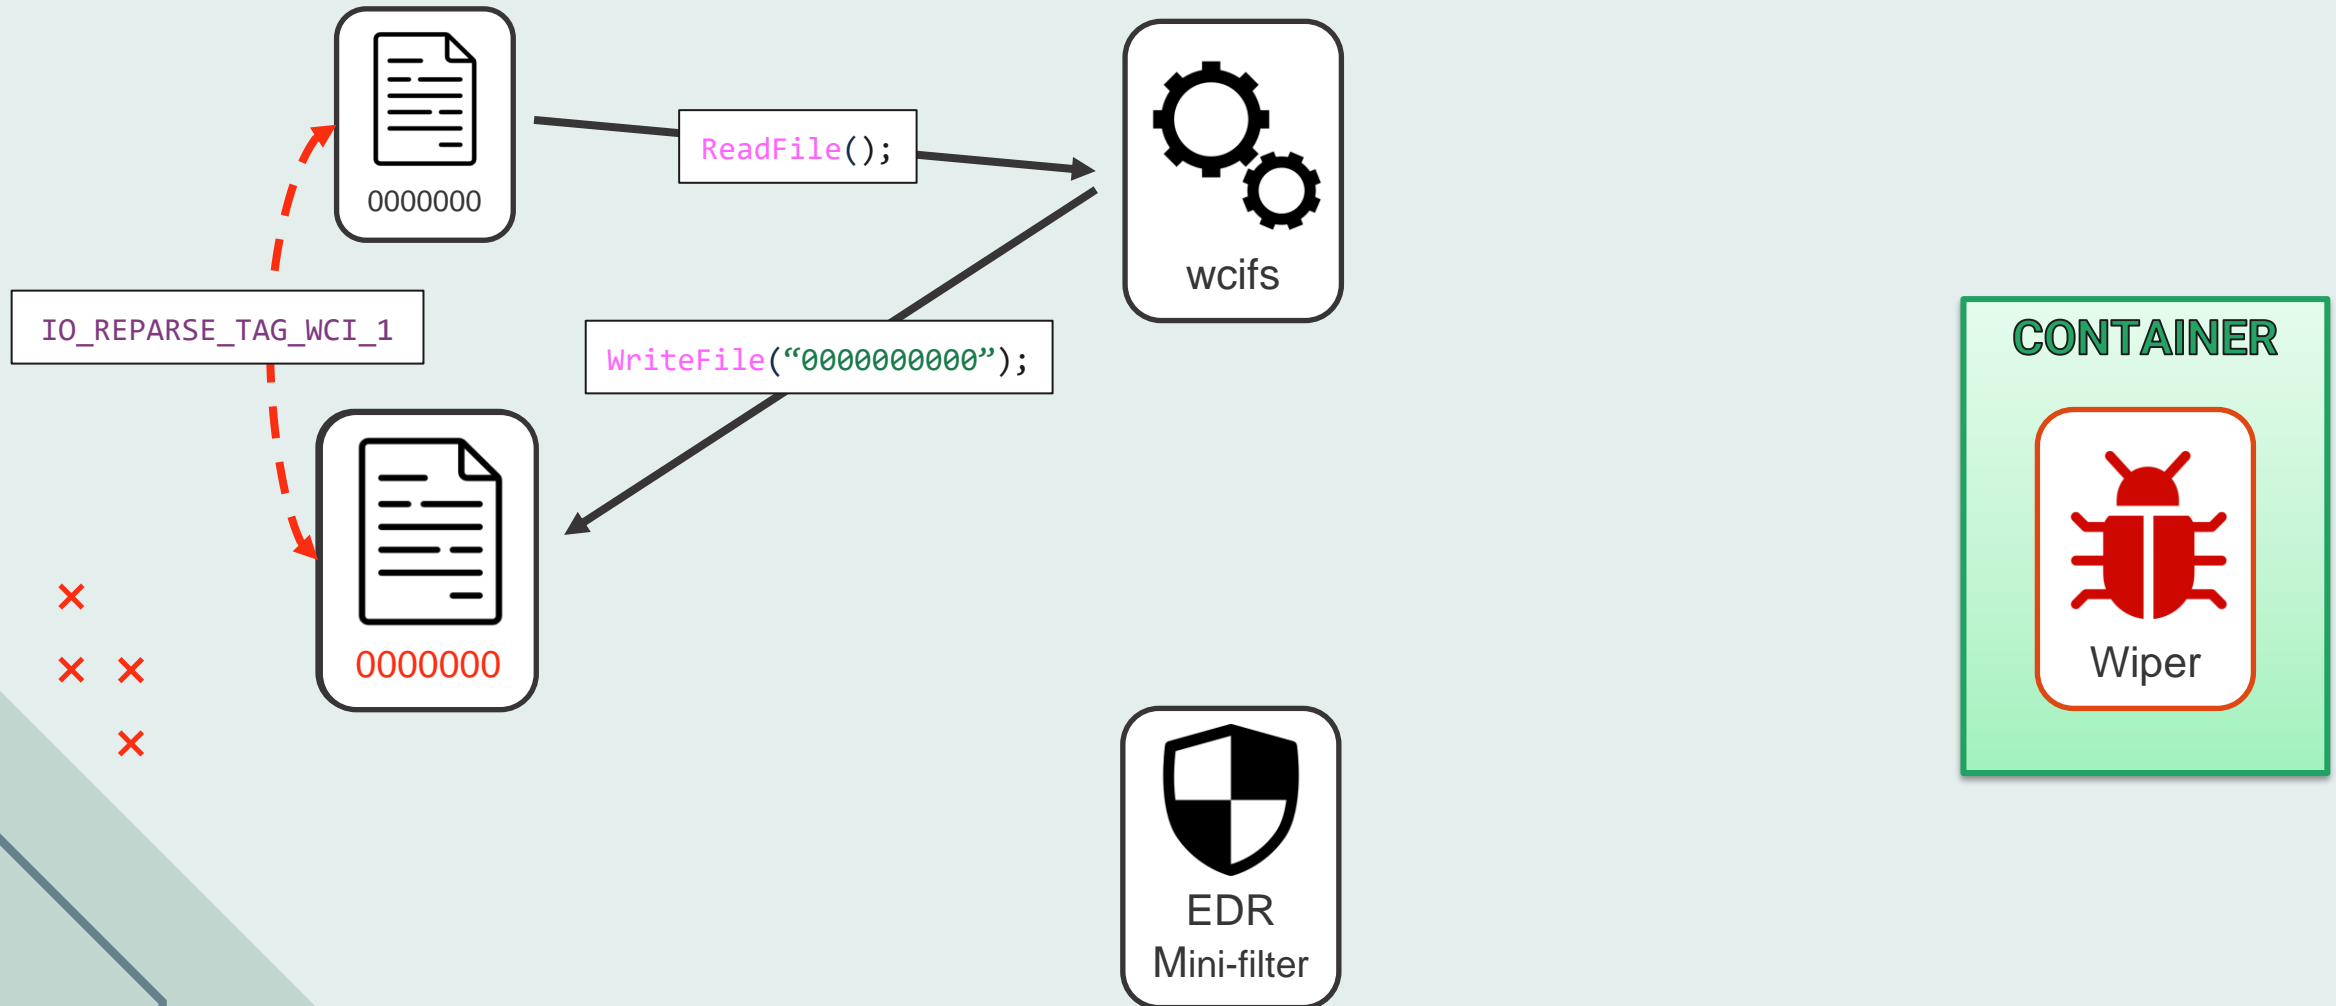
Security Products Mini-filter Bypass

Creating an Undetectable Wiper



Security Products Mini-filter Bypass

Creating an Undetectable Wiper



Security Products Mini-filter Bypass

Creating an Undetectable Ransomware

1. Traverse each file on the system and for each:

I. Read its content and encrypt in-memory

II. Create a target file and write the encrypted data to it – **will be ignored by security mini-filter because the data is written to a new file and not overriding existing content**

III. Set an `IO_REPARSE_TAG_WCI_1` reparse point on the source file that will point to the target file

2. Create a silo, assign the current process to it and register it as a fabricated container to wcifs where both source and target volumes are the main one (`\Device\HarddiskVolume3`).

×

×

×

3. Traverse each file on the system and for each:

×

I. Open the file using `CreateFile` – the wcifs driver will write the encrypted content to the original file, **the call will not trigger security mini-filter drivers callback function**

Security Products Mini-filter Bypass

EDR 102: DLP Protection – Set Read-Only Devices & Directories

- Security vendor's products can block write operations on certain directories/volumes, which can be utilized in several ways:
 - I. Organizations often determine a read-only policy for removable devices to avoid data exfiltration
 - II. Block file writes to folders containing sensitive data
- This write protection is implemented by a mini-filter driver



Security Products ETW Bypass

EDR 103: Correlating ETW Logs

- The Windows kernel acts as a log provider of a vast number of operations occurring on the system, including those on the file system
- Vendors often consume and analyze these events for any suspicious activity (usually by cross-referencing and creating an attack flow)
- When overriding a file using the `IO_REPARSE_TAG_WCI_1` tag, the read and write operations take place within a kernel work item. This will cause the ETW log to mention the system process (PID 4) as responsible for them instead of the actual process.
- Any vendor who consumes events number 15 (Read) and 16 (Write) from the Microsoft-Windows-Kernel-File provider will receive false information

×

×

×



Security Products ETW Bypass

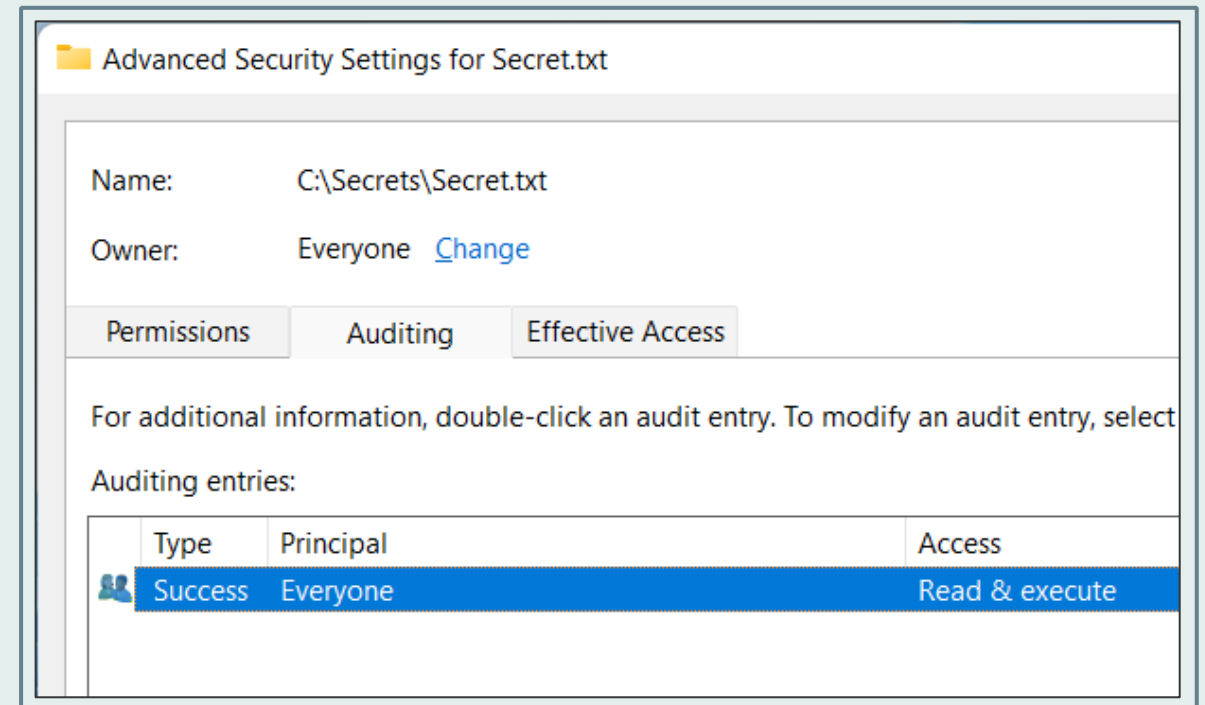
SACL Bypass

- Windows provides the option to set an auditing policy (SACL) to a file system object, which can yield a log of any I/O operation done on it
- By design, ETW-based Windows tools ignore logs that originate from the system because they should be irrelevant to a user monitoring the system (and to avoid unnecessary overhead)
- This will cause our I/O requests to be absent from the auditing logs

×

×

×




01



Overview

Windows containers
basics

02



wcifs.sys

Analyzing the main filter
driver responsible for
containers FS isolation

03



Utilization

Bypassing security
products using our
findings

04



Summary

Summary, mitigation
and future research

Summary

- The Windows containers framework provides a file system isolation solution that is implemented by reparse points and mini-filter drivers
- By reverse engineering the framework's main driver – wcifs.sys, we managed to create a fake container, insert our process into it and utilize the framework's I/O redirection mechanism to our advantage
 - I. Override files
 - II. Delete files
 - III. Copy & paste files
- **This allows us to perform file system calls that will not be detected by security vendors' mini-filters, and ETW-based products**
 - **Ransomware/Wiper protection bypass**
 - **DLP/Secured folders write bypass**
 - **ETW-based correlations bypass**

Summary

Mitigation

- `DeviceIoControl` + `FSCTL_SET_REPARSE_POINT` + `IO_REPARSE_TAG_WCI` – files with these tags should already exist on container's ghost volumes
- Check whether wcifs communication port was opened / a silo is created by a non-system process
- Check if the wcifs driver is attached to a volume while the containers feature is turned off



Summary

Further Research

- The wcifs driver is only one mini-filter among many, there might be more that can be utilized to perform system operations from the kernel itself
- It is possible to set reparse points on directories. Wcifs's symbols reference directory expansion and redirection handling
- **There are more mini-filter-based and ETW-based protections implemented by security products!**



THANK YOU!

Do you have any questions?



[@daniel_avinoam](https://twitter.com/daniel_avinoam)



<https://github.com/deepinstinct/ContainYourself>

References:

- [Playing in the \(Windows\) Sandbox - Check Point Research](#)
- [Who Contains the Containers?](#)
- <https://unit42.paloaltonetworks.com/what-i-learned-from-reverse-engineering-windows-containers/>
- [NTFS Reparse Points](#)
- [About Windows containers](#)
- Windows Kernel Programming, by Pavel Yosifovich

